

# Data Mining with R

## Table of Contents

- [Data Overview](#)
- [Summary Statistics](#)
  - [Mean](#)
  - [Median](#)
  - [All-in-One](#)
  - [Correlation](#)
- [Graphics](#)
  - [Histograms](#)
  - [Boxplots](#)
- [Near Zero Variance Predictors](#)
- [Linear Combinations](#)
- [Highly Correlated Variables](#)
- [Distribution](#)
- [Decision Tree](#)
- [Classification](#)
  - [SVM](#)
  - [KNN](#)
  - [SVM vs KNN](#)

Import libs:

```
library(caret)
library(data.table)
library(dplyr)
library(PerformanceAnalytics)
library(rpart.plot)
```

## Data Overview

Data Set Characteristics:	Number of Instances:	Attribute Characteristics:	Number of Attributes:	Associated Tasks:
Multivariate	1372	Real	5	Classification

**Dataset information:** Data were extracted from images that were taken from genuine and forged banknote-like specimens. For digitization, an industrial camera usually used for print inspection was used. The final images have 400 x 400 pixels. Due to the object lens and distance to the investigated object gray-scale pictures with a resolution of about 660 dpi were gained. Wavelet Transform tool were used to extract features from images.

### Attribute information:

1. variance of Wavelet Transformed image (type: continuous)

2. skewness of Wavelet Transformed image (type: continuous)
3. curtosis of Wavelet Transformed image (type: continuous)
4. entropy of image (type: continuous)
5. class (type: integer)

**Data source:** <https://archive.ics.uci.edu/ml/datasets/banknote+authentication>

Load data\_banknote\_authentication.txt file:

```
url = paste('https://archive.ics.uci.edu/ml/machine-learning-databases/00267/',
            'data_banknote_authentication.txt', sep='')
df = data.frame(fread(url))
names(df) = c('variance', 'skewness', 'curtosis', 'entropy', 'class')
```

Check size of df dataframe:

```
nrow(df)
```

**Output:**

```
1372
```

Show the first part of df dataframe:

```
head(df, 5)
```

**variance skewness curtosis entropy class**

```
3.62160 8.6661 -2.8073 -0.44699 0
4.54590 8.1674 -2.4586 -1.46210 0
3.86600 -2.6383 1.9242 0.10645 0
3.45660 9.5228 -4.0112 -3.59440 0
0.32924 -4.4552 4.5718 -0.98880 0
```

Show the last part of df dataframe:

```
tail(df, 5)
```

**variance skewness curtosis entropy class**

```
1368 0.40614 1.34920 -1.4501 -0.55949 1
1369 -1.38870 -4.87730 6.4774 0.34179 1
1370 -3.75030 -13.45860 17.5932 -2.77710 1
1371 -3.56370 -8.38270 12.3930 -1.28230 1
1372 -2.54190 -0.65804 2.6842 1.19520 1
```

## Summary Statistics

### Mean

```
print(noquote(paste0('Mean. Variance of Wavelet Transformed image: ',
mean(df$variance))))
print(noquote(paste0('Mean. Skewness of Wavelet Transformed image: ',
mean(df$skewness))))
print(noquote(paste0('Mean. Curtosis of Wavelet Transformed image: ',
mean(df$curtosis))))
print(noquote(paste0('Mean. Entropy of image: ', mean(df$entropy))))
```

### Output:

```
[1] Mean. Variance of Wavelet Transformed image: 0.433735257069971
[1] Mean. Skewness of Wavelet Transformed image: 1.92235312063936
[1] Mean. Curtosis of Wavelet Transformed image: 1.39762711726676
[1] Mean. Entropy of image: -1.19165652004373
```

## Median

```
print(noquote(paste0('Median. Variance of Wavelet Transformed image: ',
median(df$variance))))
print(noquote(paste0('Median. Skewness of Wavelet Transformed image: ',
median(df$skewness))))
print(noquote(paste0('Median. Curtosis of Wavelet Transformed image: ',
median(df$curtosis))))
print(noquote(paste0('Median. Entropy of image: ', median(df$entropy))))
```

### Output:

```
[1] Median. Variance of Wavelet Transformed image: 0.49618
[1] Median. Skewness of Wavelet Transformed image: 2.31965
[1] Median. Curtosis of Wavelet Transformed image: 0.61663
[1] Median. Entropy of image: -0.58665
```

## All-in-One

```
print(noquote('Summary:'))
summary(select(df, -class))
```

### Output:

```
[1] Summary:

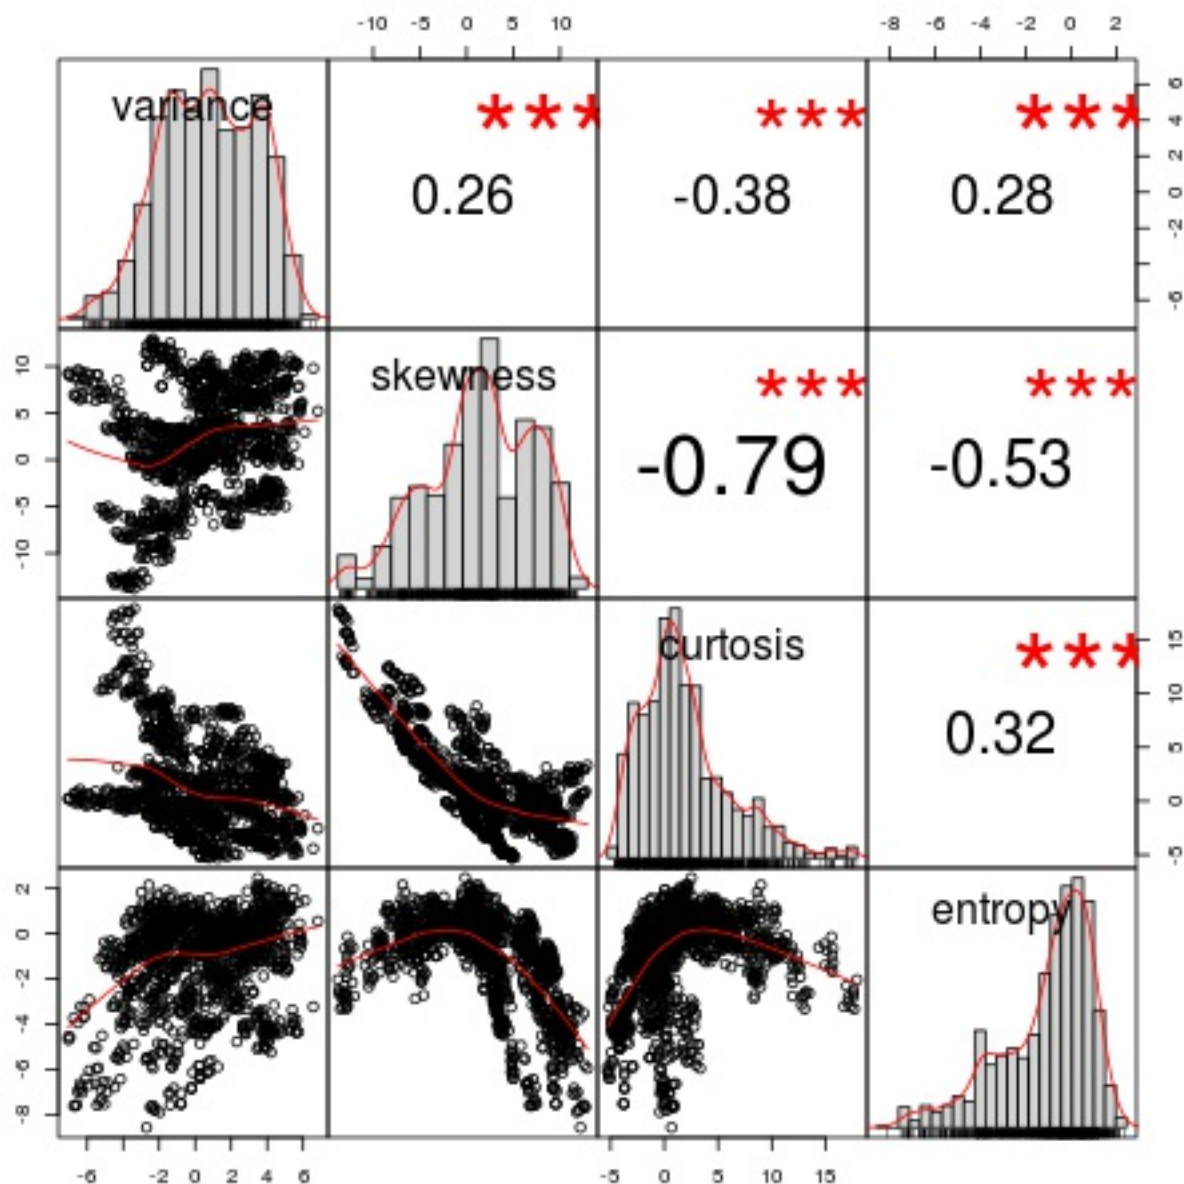
      variance      skewness      curtosis      entropy
Min.   :-7.0421  Min.   :-13.773  Min.   :-5.2861  Min.   :-8.5482
1st Qu.: -1.7730  1st Qu.: -1.708  1st Qu.: -1.5750  1st Qu.: -2.4135
Median :  0.4962  Median :  2.320  Median :  0.6166  Median : -0.5867
Mean    :  0.4337  Mean    :  1.922  Mean    :  1.3976  Mean    : -1.1917
3rd Qu.:  2.8215  3rd Qu.:  6.815  3rd Qu.:  3.1793  3rd Qu.:  0.3948
Max.    :  6.8248  Max.    : 12.952  Max.    :17.9274  Max.    :  2.4495
```

## Correlation

```
cor(df)
```

	variance	skewness	curtosis	entropy	class
<b>variance</b>	1.0000000	0.2640255	-0.3808500	0.27681670	-0.72484314
<b>skewness</b>	0.2640255	1.0000000	-0.7868952	-0.52632084	-0.44468776
<b>curtosis</b>	-0.3808500	-0.7868952	1.0000000	0.31884089	0.15588324
<b>entropy</b>	0.2768167	-0.5263208	0.3188409	1.00000000	-0.02342368
<b>class</b>	-0.7248431	-0.4446878	0.1558832	-0.02342368	1.00000000

```
chart.Correlation(select(df, -class), histogram=TRUE)
```



## Graphics


### Histograms

```

par(mfrow=c(2,2))
hist(df$variance, main='Histogram of Variance',
     xlab='Variance of Wavelet Transformed Image')
hist(df$skewness, main='Histogram of Skewness',
     xlab='Skewness of Wavelet Transformed Image')
hist(df$curtosis, main='Histogram of Curtosis',
     xlab='Curtosis of Wavelet Transformed Image')
hist(df$entropy, main='Histogram of Entropy',
     xlab='Entropy of Image')


```

**Histogram of Variance**




Variance of Wavelet Transformed Image

**Histogram of Skewness**




Skewness of Wavelet Transformed Image

**Histogram of Curtosis**



Curtosis of Wavelet Transformed Image

**Histogram of Entropy**



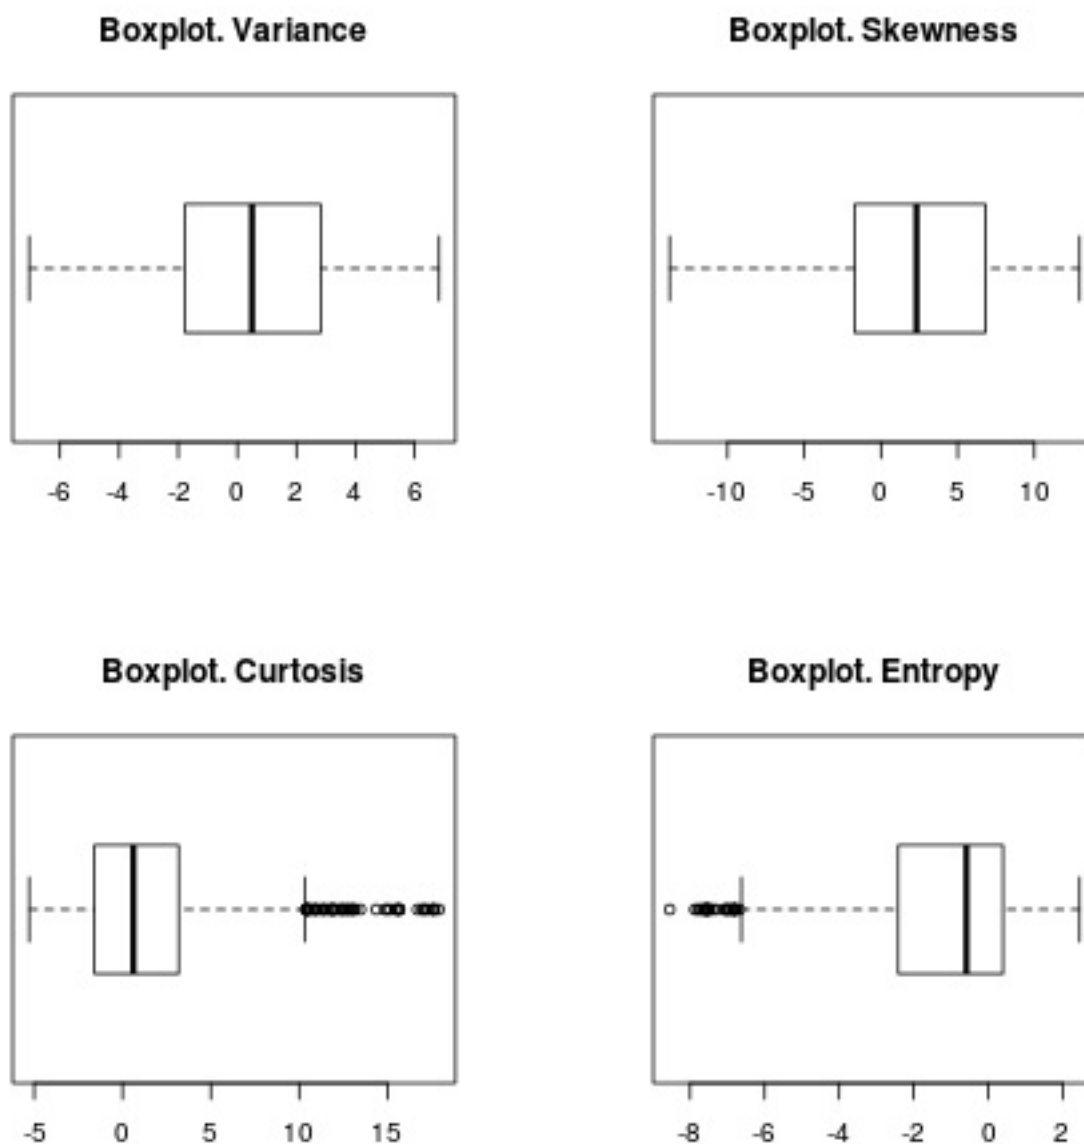
Entropy of Image

## Boxplots

```

par(mfrow=c(2,2))
boxplot(df$variance, data=df, main='Boxplot. Variance', horizontal=TRUE)
boxplot(df$skewness, data=df, main='Boxplot. Skewness', horizontal=TRUE)
boxplot(df$curtosis, data=df, main='Boxplot. Kurtosis', horizontal=TRUE)
boxplot(df$entropy, data=df, main='Boxplot. Entropy', horizontal=TRUE)

```



## Near Zero Variance Predictors

```

nearZeroVar(select(df, -class), saveMetrics=TRUE)

```

	freqRatio	percentUnique	zeroVar	nzv
<b>variance</b>	1.25	97.52187	FALSE	FALSE
<b>skewness</b>	1.20	91.54519	FALSE	FALSE
<b>curtosis</b>	1.00	92.56560	FALSE	FALSE
<b>entropy</b>	1.00	84.25656	FALSE	FALSE

## Linear Combinations

```
findLinearCombos(select(df, -class))
```

**Output:**

```
$linearCombos  
list()  
  
$remove  
NULL
```

## Highly Correlated Variables

```
df$class = as.character(ifelse(df$class=='1', 'Y', 'N'))  
df2 = select(df, -class)  
cor_matrix = cor(df2)  
print(noquote('Highly correlated variables:'))  
summary(cor_matrix[upper.tri(cor_matrix)]) # upper triangular part of a matrix
```

**Output:**

```
[1] Highly correlated variables:  
  
      Min. 1st Qu.  Median    Mean 3rd Qu.    Max.  
-0.78690 -0.48995 -0.05841 -0.13906  0.27362  0.31884
```

```
high_cor_var = findCorrelation(cor_matrix, cutoff = 0.75) # check var above  
0.75  
print(noquote(paste0('Highly correlated variables: ', names(df2)  
[high_cor_var])))
```

**Output:**

```
[1] Highly correlated variables: skewness
```

Delete highly correlated skewness column from dataframe:

```
df2 = select(df2, -skewness)
```

```
cor_matrix = cor(df2)  
print(noquote('Highly correlated variables:'))  
summary(cor_matrix[upper.tri(cor_matrix)]) # upper triangular part of a matrix  
df = cbind.data.frame(df2, class = df$class) # add class
```

**Output:**

```
[1] Highly correlated variables:  
  
      Min. 1st Qu.  Median    Mean 3rd Qu.    Max.  
-0.38085 -0.05202  0.27682  0.07160  0.29783  0.31884
```

## Distribution

```
print(noquote('Distribution:'))
table(df$class)
```

### Output:

```
[1] Distribution:
```

```
  N    Y
762 610
```

```
class_freq = data.frame(table(df$class))
names(class_freq) = c('class', 'freq')
percent_chart = cbind(class_freq,
                      percent=round((class_freq$freq/sum(class_freq$freq))*100,
1))
percent_chart
```


### class freq percent

```
N    762 55.5
Y    610 44.5
```

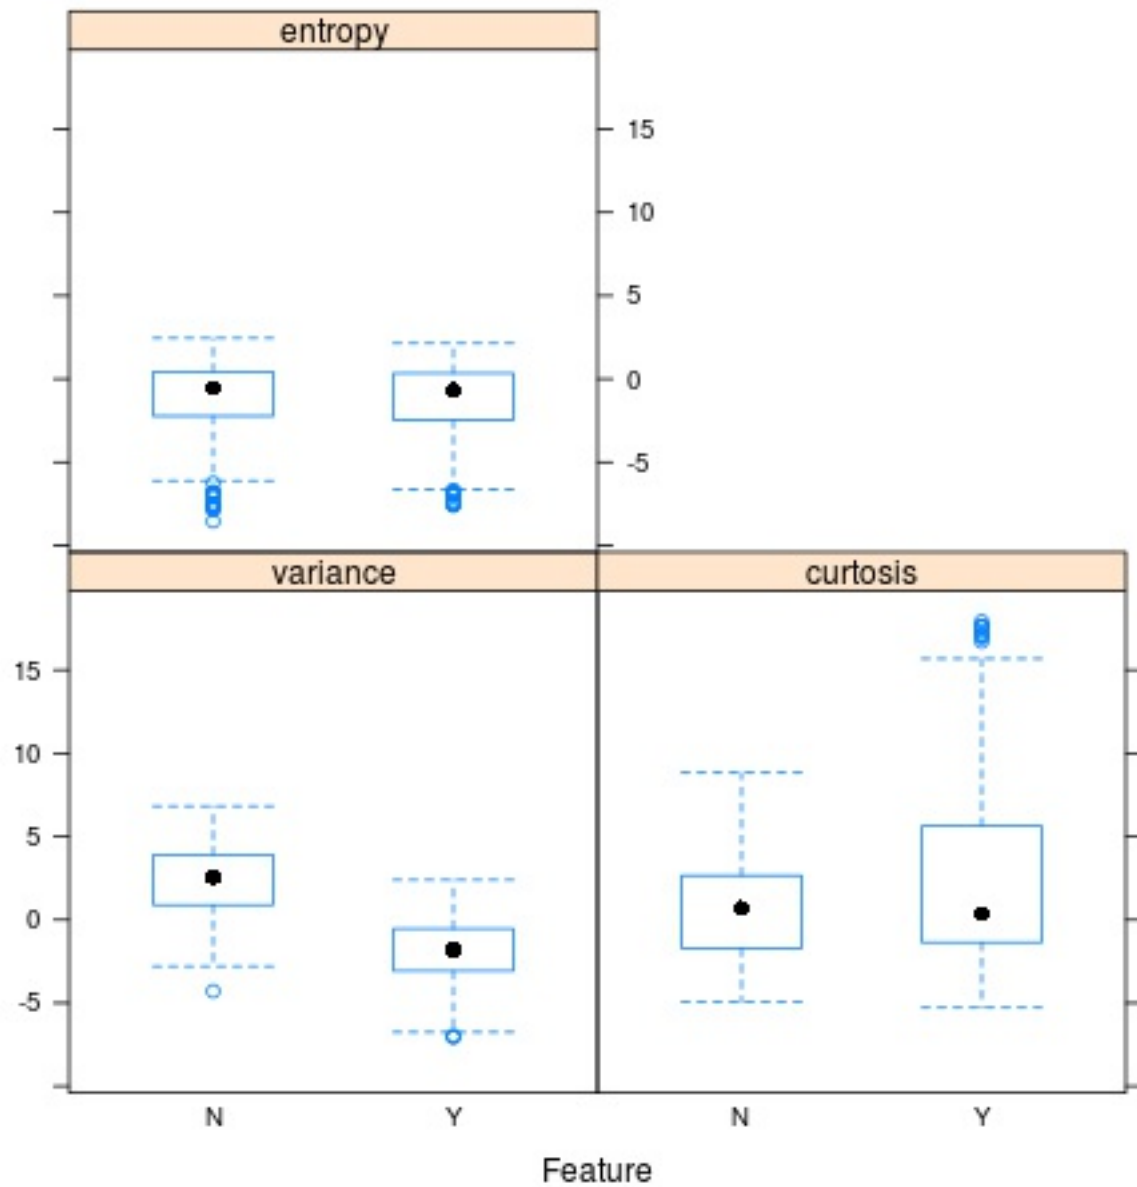
```
slices = percent_chart$percent
lbls = c('N', 'Y')
pct = round(slices/sum(slices)*100, 1)
lbls = paste(lbls, pct) # add values of pct to labels
lbls = paste(lbls, '%', sep='') # add % char to labels
pie(slices, labels=lbls, radius=1, main='Pie Chart of Distribution',
    clockwise=TRUE)
```



**Pie Chart of Distribution**

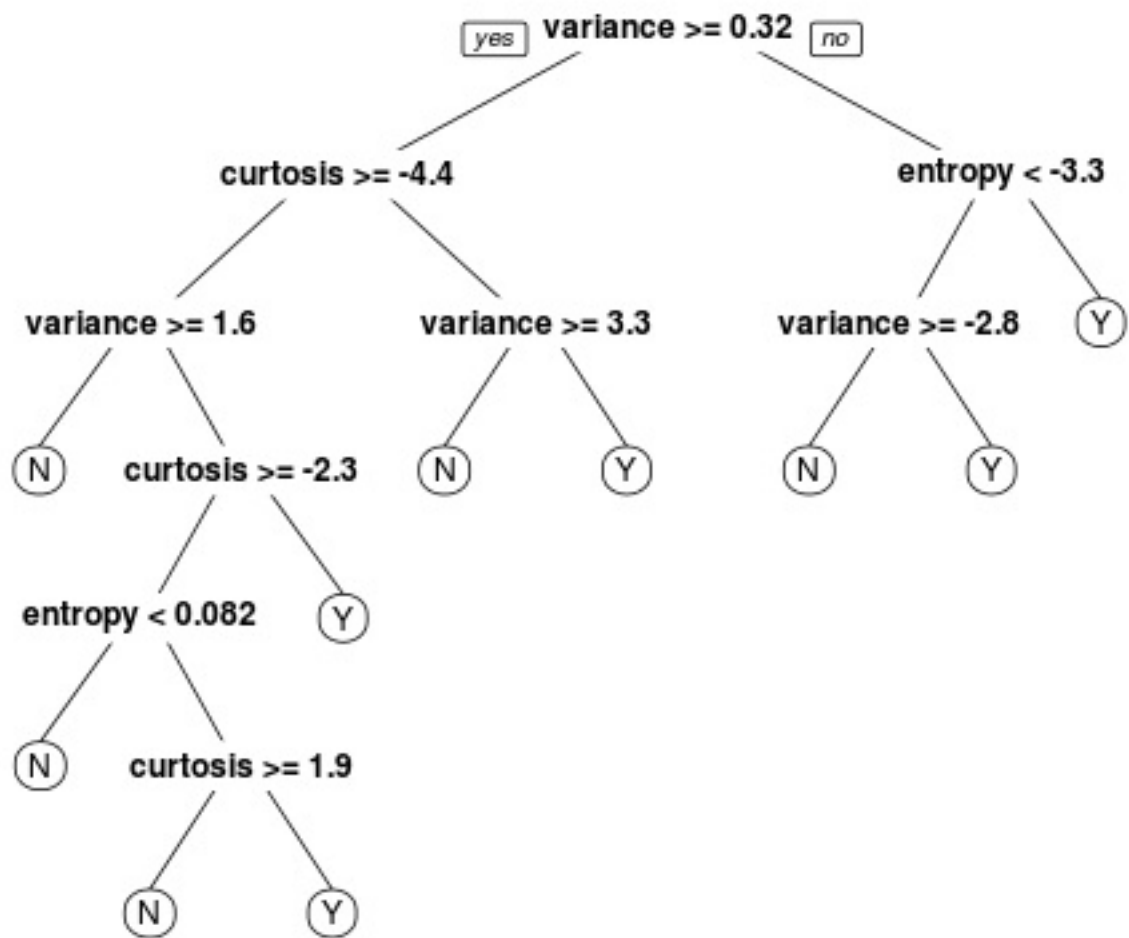


```
featurePlot(x=select(df, -class), y=df$class, plot='box')
```

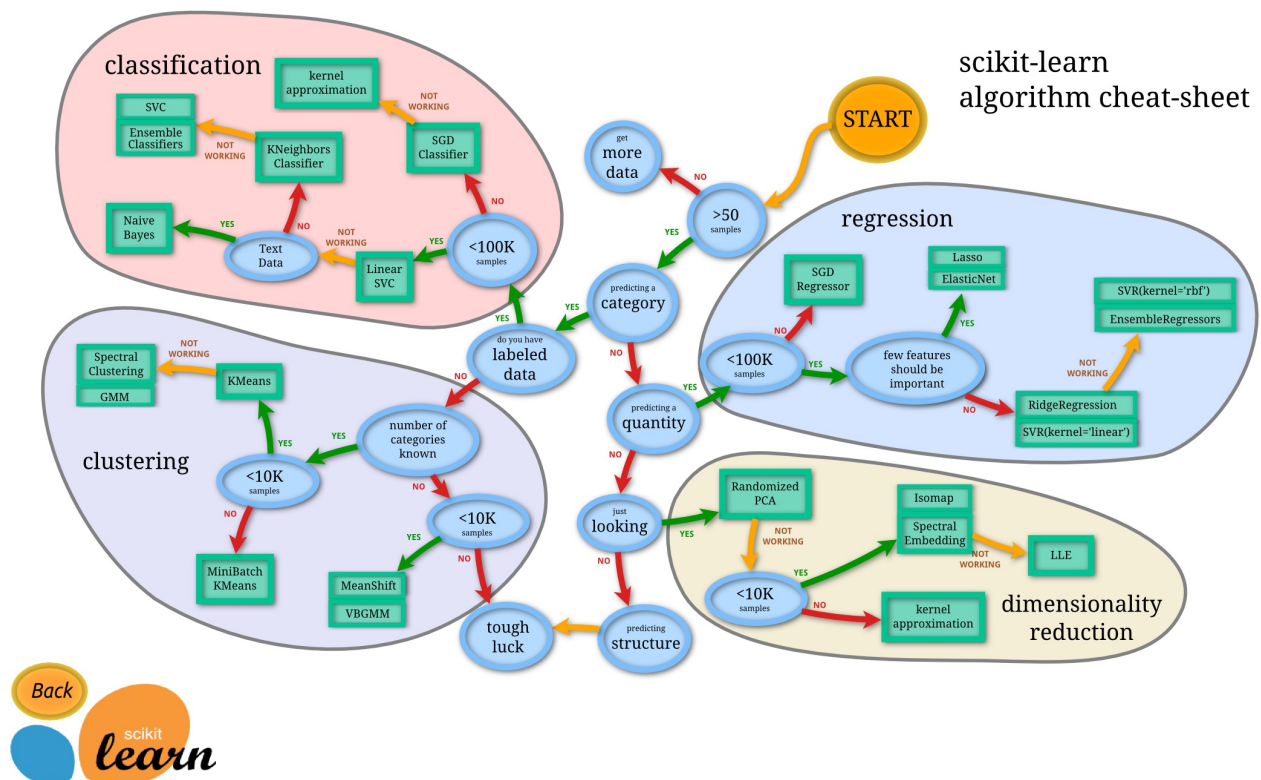


## Decision Tree

```
rmtree_set = rpart(class ~ ., df)
prp(rmtree_set)
```



**Classification**



Split the data to train and test sets:

```
train_ind = createDataPartition(df$class, p=0.7, list=FALSE)
data_train = data.frame(df[train_ind, ])
data_test = data.frame(df[-train_ind, ])
print(noquote('Train:'))
table(data_train$class)
print(noquote('Test:'))
table(data_test$class)
```

**Output:**

[1] Train:

N	Y
534	427

[1] Test:

N	Y
228	183

Choose validation method for the test of model:

```
valid_par = trainControl(method='repeatedcv', number=5, repeats=10, p=0.70,
preProc='range')
```

**SVM**

```
mod_svm = train(class ~ ., data=data_train, trControl=valid_par,
method='svmRadial')
mod_svm
```

### Output:

Support Vector Machines **with** Radial Basis Function Kernel

```
961 samples
  3 predictors
  2 classes: 'N', 'Y'
```

No pre-processing

Resampling: Cross-Validated (5 fold, repeated 10 times)

Summary **of** sample sizes: 768, 769, 770, 769, 768, 770, ...

Resampling results across tuning parameters:

C	Accuracy	Kappa
0.25	0.9785620	0.9568214
0.50	0.9807452	0.9612074
1.00	0.9815764	0.9628818

Tuning parameter '**sigma**' was held **constant at a value of 0.8345372**

Accuracy was used **to select the** optimal model **using the largest value.**

The final values used **for the** model were sigma = **0.8345372** and C = **1.**

## KNN

```
mod_knn = train(class ~. , data=data_train, trControl=valid_par, method='knn')
mod_knn
```

### Output:

k-Nearest Neighbors

```
961 samples
  3 predictors
  2 classes: 'N', 'Y'
```

No pre-processing

Resampling: Cross-Validated (5 fold, repeated 10 times)

Summary **of** sample sizes: 768, 769, 769, 769, 769, 768, ...

Resampling results across tuning parameters:

k	Accuracy	Kappa
5	0.9761700	0.9519131
7	0.9766871	0.9529962
9	0.9748131	0.9492504

Accuracy was used **to select the** optimal model **using the largest value.**

The final **value** used **for the** model was k = **7.**

Show summary:

```
print(noquote('Summary:'))
mod_results = resamples(list(SVM=mod_svm, KNN=mod_knn))
summary(mod_results)
```

### Output:

```
[1] Summary:

Call:
summary.resamples(object = mod_results)


Models: SVM, KNN
Number of resamples: 50

Accuracy
      Min.   1st Qu.   Median     Mean   3rd Qu.   Max. NA's
SVM 0.9633508 0.9740933 0.9843750 0.9815764 0.9882606 1.0000000    0
KNN 0.9581152 0.9687905 0.9740933 0.9766871 0.9843750 0.9947917    0

Kappa
      Min.   1st Qu.   Median     Mean   3rd Qu.   Max. NA's
SVM 0.9264050 0.9478632 0.9684487 0.9628818 0.9762997 1.0000000    0
KNN 0.9157941 0.9372235 0.9476255 0.9529962 0.9684487 0.9894575    0
```

### SVM vs KNN

```
bwplot(mod_results, scales=list(x=list(relation='free'),
y=list(relation='free')))
```



Test models:

```
test = select(data_test, -class)
test_sum = data_test$class
mod_predict_svm = predict(mod_svm, test)
print(noquote('SMV:'))
confusionMatrix(mod_predict_svm, test_sum)
mod_predict_knn = predict(mod_knn, test)
print(noquote('KNN:'))
confusionMatrix(mod_predict_knn, test_sum)
```

**Output:**

```
[1] SMV:
```

Confusion Matrix **and** Statistics

```

      Reference
Prediction  N    Y
      N 222    1
      Y   6 182

      Accuracy : 0.983
      95% CI : (0.9652, 0.9931)
      No Information Rate : 0.5547
      P-Value [Acc > NIR] : <2e-16

      Kappa : 0.9656
      McNemar's Test P-Value : 0.1306

      Sensitivity : 0.9737
      Specificity : 0.9945
      Pos Pred Value : 0.9955
      Neg Pred Value : 0.9681
      Prevalence : 0.5547
      Detection Rate : 0.5401
      Detection Prevalence : 0.5426
      Balanced Accuracy : 0.9841

      'Positive' Class : N

```

[1] KNN:

Confusion Matrix **and** Statistics

```

      Reference
Prediction  N    Y
      N 223    2
      Y   5 181

      Accuracy : 0.983
      95% CI : (0.9652, 0.9931)
      No Information Rate : 0.5547
      P-Value [Acc > NIR] : <2e-16

      Kappa : 0.9656
      McNemar's Test P-Value : 0.4497

      Sensitivity : 0.9781
      Specificity : 0.9891
      Pos Pred Value : 0.9911
      Neg Pred Value : 0.9731
      Prevalence : 0.5547
      Detection Rate : 0.5426
      Detection Prevalence : 0.5474
      Balanced Accuracy : 0.9836

      'Positive' Class : N

```